

Synchronized state:

For the synchronized state the swap function is made synchronized. This creates a monitor acquire event while a thread enters a function and a monitor release when a thread is done with the executing the function. In short, no thread B can interrupt a thread A if thread A is executing the synchronized function, swap(int I, int j). 100% reliable but for large values of # of threads tends to get very slow.

Unsynchronized state:

For unsynchronized state, I just removed the synchronized keyword and kept the rest of the code the same. There are ample places for the data race condition to occur, one such occurrence being a thread being interrupted when value[i] and value[j] is being read. If the interrupting thread now updates the value of value[i] and value[j] then we have a data race. Not 100% reliable but very fast performance.

GetNSet State:

For this state implementation, I used the java.util.concurrent.Atomic library. All the read/write operations performed are atomic and the byte array was converted to an atomic integer array. This will always be DRF since all the portions which induce a DR are made atomic.

BetterSafe State:

For this I looked into CAS synchronization. I had a flag (**AtomicBoolean type**) called “locked” to indicate whether the data race inducing portion (the read and write of values into the byte array) of the code is locked or not. I have pasted a code snippet to illustrate the point.

```
public boolean swap(int i, int j) {  
    if (locked.equals(false)){//check the flag to see if the volatile resource is locked or not  
        locked.compareAndSet(false, true); //LOCK this part of the code for a thread to perform read/write operation  
        if (value[i] <= 0 || value[j] >= maxval) {  
            locked.compareAndSet(true, false); //UNLOCK after READ  
            return false;  
        }  
        value[i]--;  
        value[j]++;  
        locked.compareAndSet(true, false);// or UNLOCK after WRITE  
    }  
    return true;  
}
```

This implementation of state is also DRF since the atomic boolean flag, locked, acts as a lock on the read/write operations on value[i] and value[j]. If a thread is interrupted, it can only be interrupted after locking the volatile part hence making it inaccessible for the other threads. **This is very fast and DRF.**

I also considered other approaches like making the individual elements of the array atomic and then performing read/write operations on them, but even though that remains DRF, I found it to be slower than the current implementation. I also tried using **Reentrant Locks** but still found them slower than this approach.


This approach is 100% reliable and fast performance.

BetterSorry State:

For BetterSorry implementation of state, I mapped the read operations on to local variables, but the write operations were done non-atomically. This also is as fast as BetterSafe (only occasionally being faster than the BetterSafe implementation). *BetterSorry has a data race condition if the values value[i] and value[j] change and a thread is interrupted after the old values of value[i] and value[j] are copied to the local variables.*

Performance Results:

of threads VS ns /transition



As it can be seen from the graph above, Unsynchronized and BetterSorry are not DRF. BetterSafe has a much better performance than the other models which are DRF.

BetterSafe and BetterSorry (considering the fact that GDI prefers speed to efficiency) seem to be better for GDIs specifications

On Reliability:

Synchronized>=GetNSet>=BetterSafe>BetterSorry>UnSynchronized

On Performance:

UnSynchronized>BetterSorry>=BetterSafe>Synchronized>=GetNSet

Results were obtained on the following hardware:

4GB RAM